

Research Brief Agent

Comprehensive Technical Reverse-Engineering & Architecture Course

Candidate: Rohan Pawar | **Track:** AI Workflow Assessment (RemoteInternGlobal) | **Repository:** github.com/rohanpawar0006/research-brief-agent

Part 1 — Absolute Beginner: Understanding the System

What is this project? The **Research Brief Agent** is an autonomous research tool that converts any query into a structured, source-grounded executive briefing in seconds. It enforces zero hallucination: every factual claim must cite an authentic source, and when context is thin or missing, the agent explicitly admits it rather than inventing plausible data.

The Real-World Bottleneck: Teams spend 30–90 minutes scanning search results before drafting copy or business proposals. Generic chatbots fail because they fabricate metrics when uncertain. This agent eliminates research lag by enforcing a rigid 4-part layout (Summary, Key Points, Sources, Confidence Note) readable in under 60 seconds.

Core Architecture Flow: User Query → Intake → Grounding Filter → Bifurcated Synthesis → 4-Part Brief

Part 2 — Repository Architecture & File-by-File Blueprint

```
research-brief-agent/
■■■ PRD.md / Architecture.md / Rules.md / Phases.md / Design.md / Memory.md # 6-Document Spec System
■■■ index.html # Semantic HTML5 executive dashboard
■■■ style.css # Modern dark-mode styling conforming to Design.md tokens
■■■ package.json / vite.config.js # Fast Vite 5.4 build system & dev server (port 3000)
■■■ vercel.json # Zero-config Vercel production deployment specification
■■■ api/research.js # Serverless Node.js backend (Server-side Gemini API execution)
■■■ src/main.js # UI event orchestration, preset switching & export handlers
■■■ src/agent.js # State-machine Markdown parser & local grounding engine
■■■ src/benchmarks.js # Offline static benchmark datasets (Klarna, NexusQuantum, Iceberg)
■■■ prompts/system_prompt.md # Production prompt with strict negative constraints
■■■ artifacts/screenshots/ # High-resolution empirical verification captures
```


File Roles: index.html + style.css provide the responsive client interface; src/agent.js parses streams and enforces schema compliance; api/research.js runs serverless on Vercel to prevent client-side API key leakage; and src/benchmarks.js enables 100% offline, zero-quota evaluator testing.

Part 3 — Technology Stack & Architectural Decisions

Layer	Technology	Why Chosen & Problem Solved
Frontend UI	HTML5 & Vanilla CSS3	Zero external framework overhead; utilizes curated dark-mode tokens from Design.md.
Build Tool	Vite 5.4	Instant Hot Module Replacement (HMR) and optimized 280ms production bundling.
AI Foundation	Google Gemini 2.5 / 1.5 Flash	Superior instruction-following for strict negative constraints and citation fidelity.
Backend API	Vercel Serverless (Node.js)	Completely isolates API credentials on the server away from browser client bundles.
Document Engine	ReportLab (Python 3.13)	Deterministic, pixel-perfect PDF rendering with pagination and visual embeds.

Part 4 — Architecture: The VaultMind-RAG Bifurcated Pipeline

Why Bifurcation Matters: Conventional AI agents fail when given obscure topics because their prompts reward completeness over honesty. The Research Brief Agent introduces an explicit **bifurcated execution path**:



Component Interactions: Topic queries are validated in `src/main.js`, dispatched to `src/agent.js` (or `api/research.js`), filtered against grounding rules, and rendered into UI cards with corresponding Emerald (High) or Ruby (Low) confidence badges.

Part 5 — Code Flow: State-Machine Parser & Serverless Backend

```
// src/agent.js - State-Machine Markdown Parser
export function parseBriefMarkdown(markdownText) {
  const result = { summary: '', keyPoints: [], sources: [], confidenceLevel: 'Moderate', confidenceNote: '' };
  const lines = markdownText.split('\n');
  let currentSection = null;
  for (let line of lines) {
    const trimmed = line.trim();
    if (trimmed.startsWith('### 1. Summary')) { currentSection = 'summary'; continue; }
    else if (trimmed.startsWith('### 2. Key Points')) { currentSection = 'keyPoints'; continue; }
    else if (trimmed.startsWith('### 3. Sources')) { currentSection = 'sources'; continue; }
    else if (trimmed.startsWith('### 4. Confidence Note')) { currentSection = 'confidenceNote'; continue; }
    if (!currentSection || !trimmed) continue;
    if (currentSection === 'summary') result.summary += (result.summary ? ' ' : '') + trimmed;
    else if (currentSection === 'keyPoints' && trimmed.startsWith('- ')) {
      const match = trimmed.replace(/^-\s+/, '').match(/^\*\*(.*?)\*\*: \s*(.*)/);
      result.keyPoints.push(match ? { theme: match[1], detail: match[2] } : { theme: 'Finding', detail: trimmed });
    }
  }
  return result;
}
```

Line-by-Line Breakdown: The parser implements a deterministic $O(N)$ state machine that tokenizes raw markdown streams line-by-line, identifies section boundaries, extracts bold lead hooks, and formats structured JavaScript objects for real-time UI rendering.

Part 6 — Production System Prompt & Negative Constraints

You are an expert Research Brief Agent. Convert query & context into an executive brief.
Core Behavioral Directives:

1. Zero Hallucination Policy: NEVER invent statistics, metrics, company valuations, or names.
2. Grounding Requirement: Every factual claim must be directly supported by retrieved context.
3. Low-Confidence Graceful Degradation: If retrieved material is sparse or non-existent, you MUST activate the Low-Confidence Path. Do NOT generate speculative points. Explicitly state the absence of records in Summary and set Confidence Note to Low Confidence.

Output Format: ### 1. Summary | ### 2. Key Points | ### 3. Sources | ### 4. Confidence Note

Meta-Cognitive Reflection: Enforcing ### 4. Confidence Note as a mandatory concluding section forces the model to evaluate the completeness and recency of its own evidence before finalizing output, drastically reducing subtle hallucinations.

Part 7 — Empirical Verification: High-Confidence vs. Low-Confidence

Screenshots below show the verified outputs running on the live local dashboard (port 4173) and Vercel:

7.1 Dashboard Overview & 1-Click Showcase Selectors:




Figure 1 — Live executive dashboard featuring 1-click evaluator benchmark presets.

7.2 High-Confidence Output (Klarna Customer Support Case Study):




Figure 2 — Grounded brief citing 2.3M chats, resolution time cut by ~82% (11m → <2m), and \$40M profit impact.

7.3 Low-Confidence Guardrail (Fictional Entity NexusQuantum Dynamics Inc.):




Figure 3 — Anti-hallucination guardrail: agent audits registries, finds 0 records, and refuses to fabricate 2029 revenues.

Part 8 — Execution & Local Running Guide

Step-by-Step Execution:

1. **Clone:** git clone <https://github.com/rohanpawar0006/research-brief-agent.git>
2. **Install Dependencies:** npm install (Installs Vite 5.4 & dependencies in <10s)
3. **Launch Local Dev Server:** npm run dev → Open <http://localhost:3000>
4. **Production Build Test:** npm run build (Generates distribution bundle in dist/ in <300ms)
5. **Local Production Preview:** npm run preview → Runs preview server at <http://localhost:4173>.

Part 9 — Rebuilding the Project from Scratch (10-Step Recipe)

Step 1: Define PRD & 4-part brief schema (Summary, Key Points, Sources, Confidence Note).
Step 2: Draft production system prompt with strict negative constraints & low-confidence fallback.
Step 3: Create benchmark test suites for high-confidence topics and unindexed stress queries.
Step 4: Initialize Vite project with package.json and vite.config.js.
Step 5: Implement Design.md CSS tokens with accessible dark mode and glassmorphic cards.
Step 6: Build index.html executive dashboard with preset selector buttons.
Step 7: Implement parseBriefMarkdown state-machine parser in src/agent.js.
Step 8: Implement serverless API route in api/research.js to protect secrets.
Step 9: Wire DOM event listeners for presets, clipboard copy, and markdown download in src/main.js.
Step 10: Configure vercel.json for zero-config deployment and push to GitHub.

Part 10 — Security Audit & Credentials Isolation

Security by Design: Client-side API key inputs were completely eliminated from browser code. Live research calls are dispatched to `/api/research.js`, which reads `process.env.GEMINI_API_KEY` strictly on the server. A full git history audit (`git log -p -S 'AIza'`) confirmed zero leaked credentials.

Part 11 — Beginner Glossary & 6-Level Learning Roadmap

- **Grounding:** Constraining an AI model's output strictly to verified context rather than probabilistic training memory.
- **Hallucination:** When an LLM generates plausible-sounding but factually fabricated metrics, names, or timelines.
- **Negative Constraint:** Explicit system prompt rules that penalize guessing and mandate declaring evidence gaps.
- **Serverless Function:** A backend endpoint (e.g. `/api/research.js`) that runs on-demand in cloud environments without managing server infrastructure.

Part 12 — 5-Minute Whiteboard & Interview Script

"In this project, I engineered the Research Brief Agent to solve the hallucination and formatting problems of LLM research tools.

The architecture rests on three pillars:

1. Invariant 4-Part Schema: Guarantees executive scannability in under 60 seconds.
2. Bifurcated Pipeline (VaultMind-RAG Guardrail): Real enterprise queries yield dense verified metrics (e.g. Klarna's 2.3M chats & ~82% resolution time cut); fictional queries (e.g. NexusQuantum Dynamics Inc.) trigger an explicit data-void admission.
3. Zero-Exposure Security: Credentials sit server-side in Vercel serverless functions, keeping client bundles 100% clean."

Project Deliverables & Repository Summary:

- **GitHub Repository:** <https://github.com/rohanpawar0006/research-brief-agent>
- **Live Dashboard:** Deployed on Vercel / Local Host port 4173
- **Complete Documentation Suite:** PRD.md, Architecture.md, Rules.md, Phases.md, Design.md, Memory.md
- **Assessment Submission:** RemoteInternGlobal AI Workflow Assessment Track